



CAPSTONE PROJECT

ANALYZING SAUDI STOCK EXCHANGE (TADAWUL) WITH MACHINE LEARNING

DAQEST GROUP

CONTENT

- **Introduction**
- **Data Review**
- **Dataset preprocessing**
- **Exploratory Data Analysis (EDA)**
- **Dashboards**
- **Machine Learning Models**
 - **Regression Model**
 - **Classification Model**
- **Results**
- **Future Work**

INTRODUCTION



- Saudi's 2030 Vision and Tadawul
- Business problem
- Project goal
 - Predicts the close price
 - Predict the stock change

INTRODUCTION



Data Review

TADAWAL DATASET

- This is the data of Saudi stock market companies from 2001-12-31 until 2020-04-16.
- Collected from Saudi Stock Exchange (Tadawul) website.
- This dataset has 14 columns and 593819 rows.

TADAWAL DATASET

Each row in the database represents the price of a specific stock at a specific date:

(symbol (Integer	The symbol or the reference number of the company.
(name(String	Name of the company.
(trading_name (String	The trading name of the company
(sectoer (String	The sector in which the company operates.
(date (Date	The date of the stock price.
(open (Decimal	The opening price.
(high (Decimal	The highest price of the stock at that day.
(low (Decimal	The lowest price of the stock at that day.
(close (Decimal	The closing price.
(change (Decimal	The change in price from the last day.
(perc_Change (Decimal	The percentage of the change.
volume_traded (Decimal)	The volume of the trades for the day.
value_traded (Decimal)	The value of the trades for the day.
(no_trades (Decimal	The number of trades for the day.

DATA REVIEW



Preprocessing Dataset

Presented By : Jawhara Almulhim

TADAWAL DATASET

- Data set shape

```
# Check the shape of dataset  
tadawul_stuks.shape
```

```
(593819, 14)
```

- Change the column Name

```
change_names = {'name': 'Company_name', 'sectoer': 'sector', 'open': 'open_price', 'high': 'high_price', 'low': 'low_price',  
                'close': 'close_price', 'trading_name': 'trading_name',  
                'volume_traded': 'volume_traded', 'no_trades': 'num_trades'}  
  
# Place them on our dataframe  
tadawul_stuks.rename(columns=change_names, inplace=True)
```


DATA PREPROCESSING

- Add year, month and day columns

```
tadawul_stuks["Year"] = pd.DatetimeIndex(tadawul_stuks['date']).year  
tadawul_stuks["month"] = pd.DatetimeIndex(tadawul_stuks['date']).month_name()  
tadawul_stuks["day"] = pd.DatetimeIndex(tadawul_stuks['date']).day_name()
```

- Drop unnecessary columns

```
#remove unnecessary column symbol, company_name  
tadawul_stuks.drop('symbol', inplace=True, axis=1)  
tadawul_stuks.drop('Company_name', inplace=True, axis=1)
```

- Add the categorical target column

```
tadawul_stuks.loc[tadawul_stuks['perc_Change'] > 0.00 , "Change_category"] = 'Good Change'  
tadawul_stuks.loc[tadawul_stuks['perc_Change'] <= -0.00 , "Change_category"] = 'Bad Change'  
tadawul_stuks.loc[tadawul_stuks['perc_Change'] == 0.00 , "Change_category"] = 'Stable'
```

DATA CLEANING

- Check the null values

```
tadawul_stuks.isnull().sum()
```

symbol	0
Company_name	0
trading_name	0
sector	0
date	0
open_price	6455
high_price	6697
low_price	6697
close_price	0
change	0
perc_Change	0
volume_traded	0
value_traded	0
num_trades	7691
Year	0
month	0
day	0
Change_category	0
dtype: int64	

- Handle the null values

```
tadawul_stuks = tadawul_stuks.dropna()
```

DATA CLEANING

DATA PREPROCESSING

- Check for null values again

```
# Check for null values again  
tadawul_stuks.isnull().sum()
```

```
trading_name      0  
sector            0  
date              0  
open_price        0  
high_price        0  
low_price         0  
close_price       0  
change            0  
perc_Change       0  
volume_traded     0  
value_traded      0  
num_trades        0  
Year              0  
month             0  
day               0  
Change_category   0  
dtype: int64
```

- Check the duplicate data

```
tadawul_stuks.duplicated().any()
```

```
False
```

DATA CLEANING


- Change the name of the sector to make it short before plotting

```
tadawul_stuks['sector'] = tadawul_stuks['sector'].replace('Information Technology', 'IT')
```

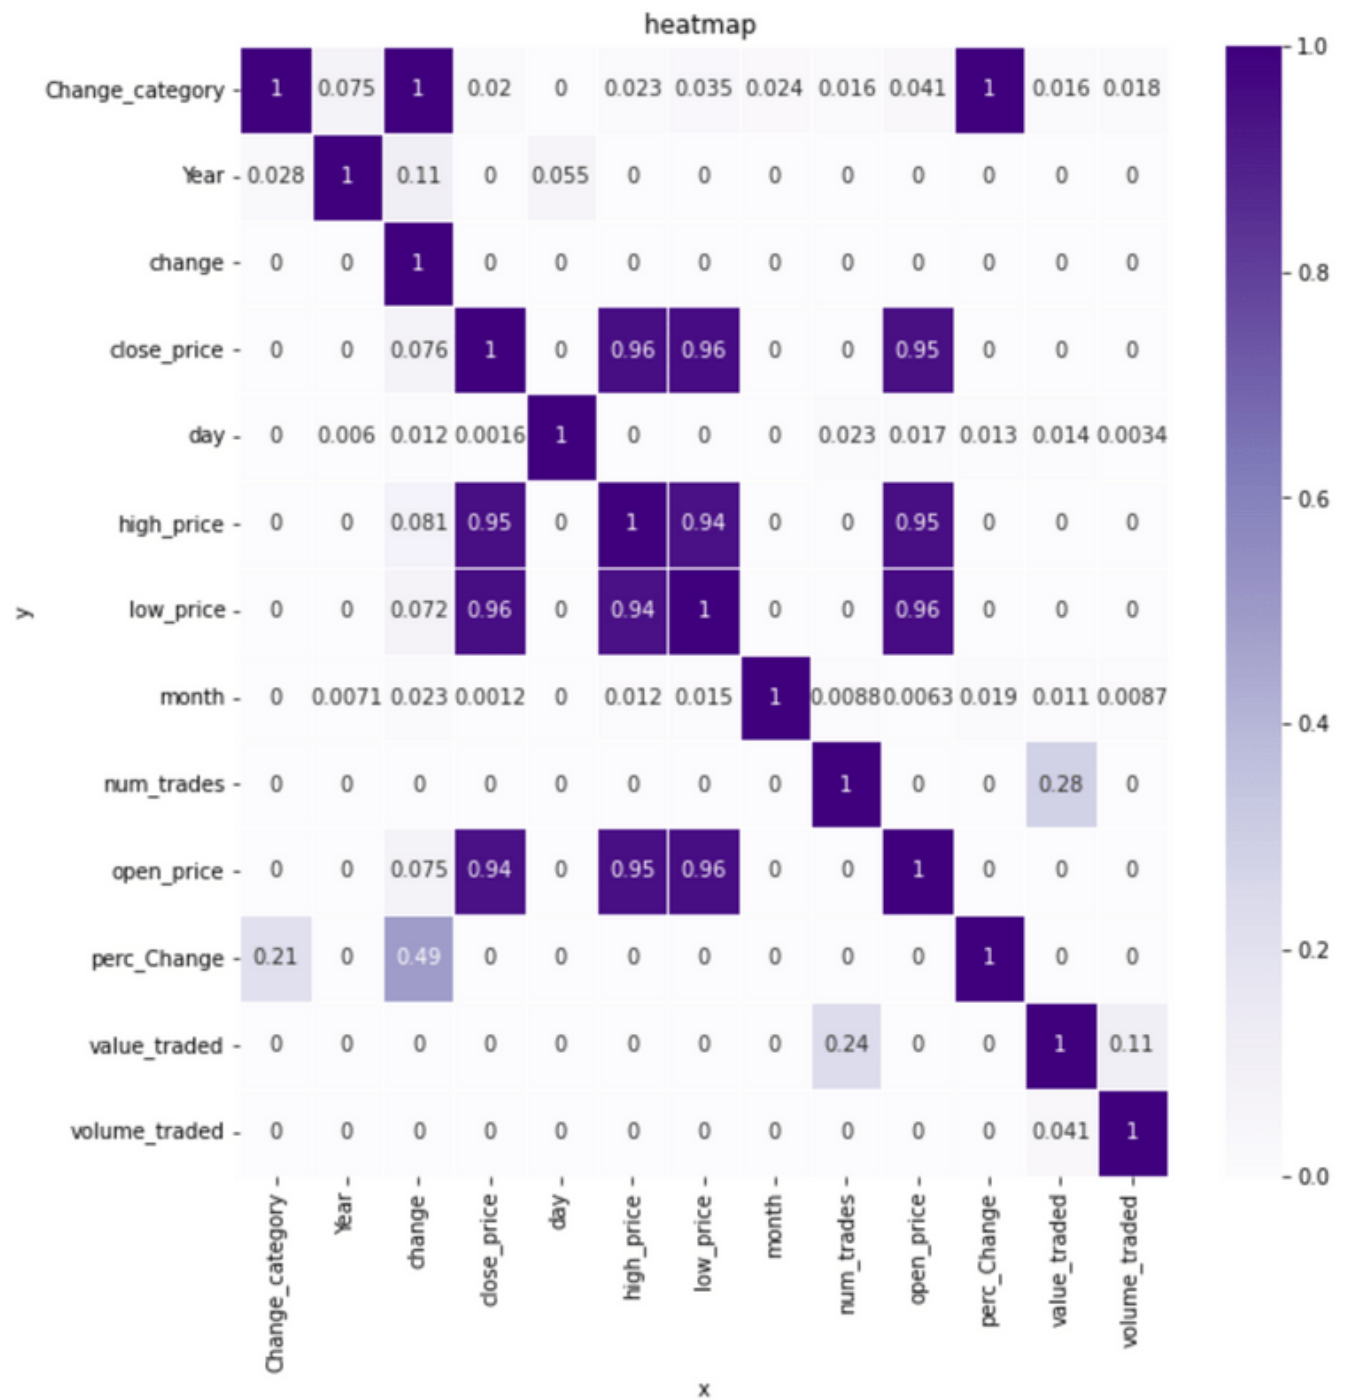
- Check the shape and size after preprocessing

```
# Check the shape of dataset  
tadawul_stuks.shape
```

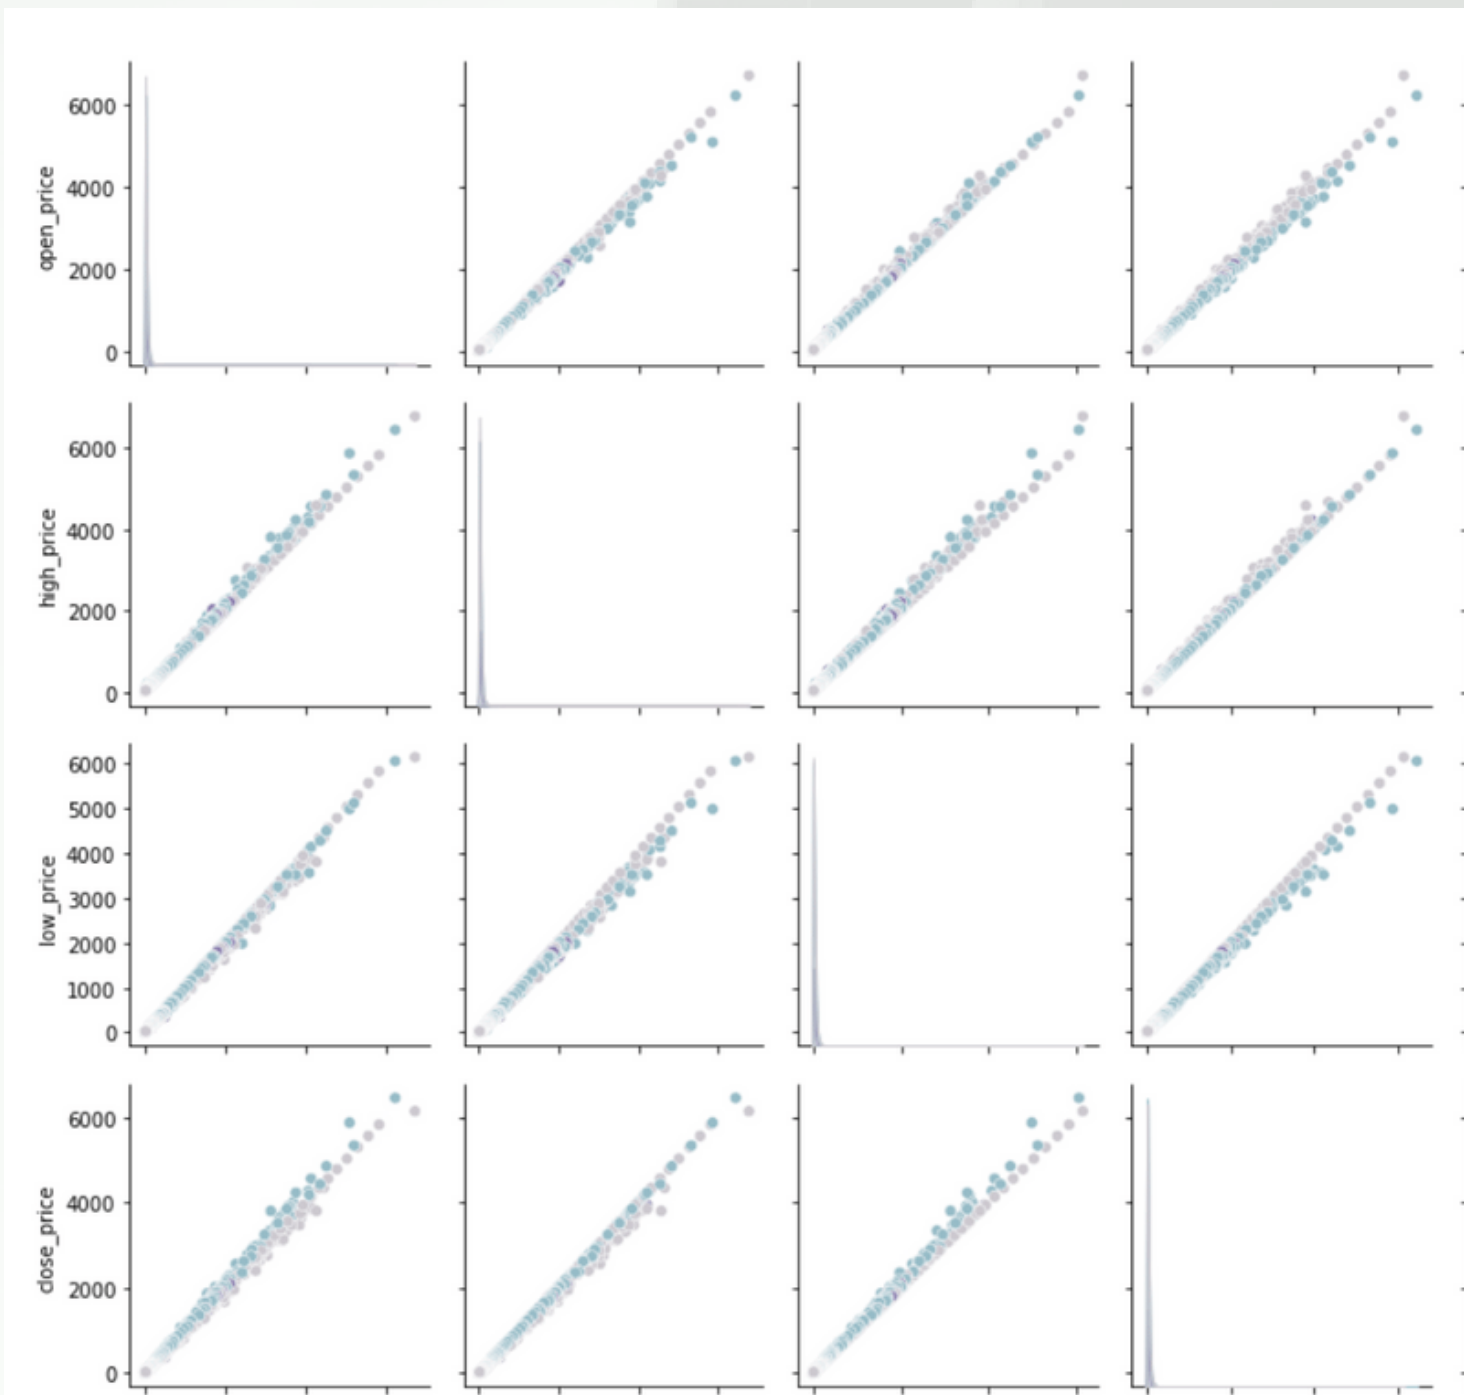
```
(579431, 16)
```

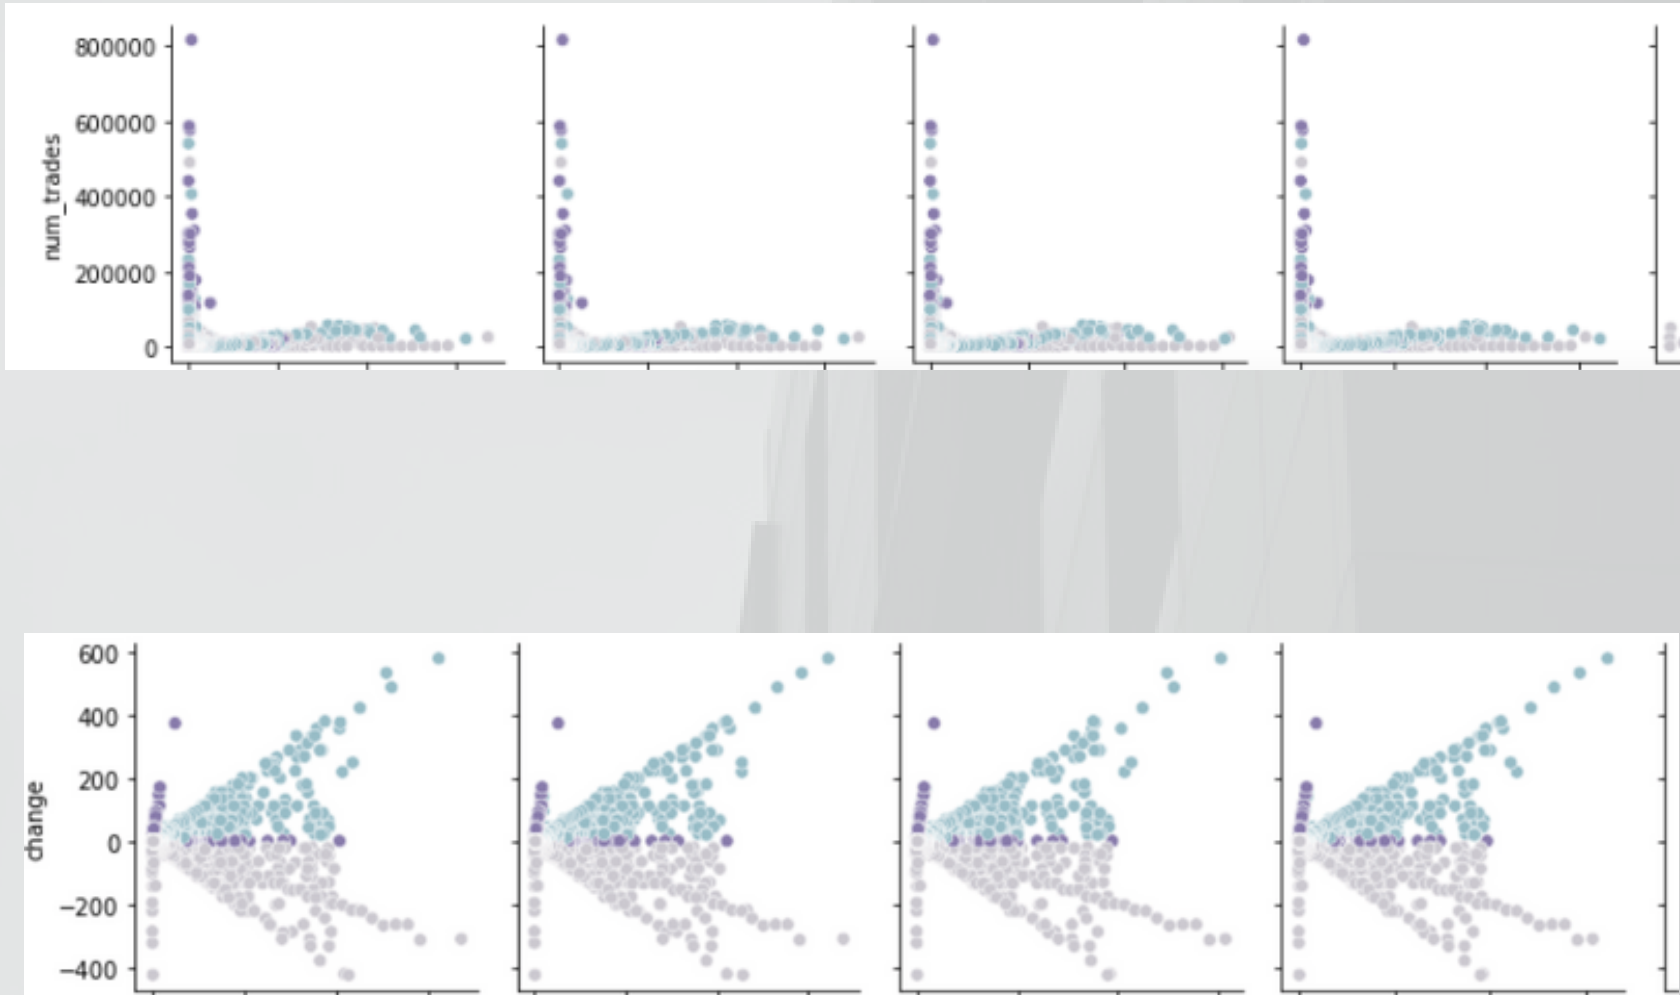



Exploratory Data Analysis (EDA)

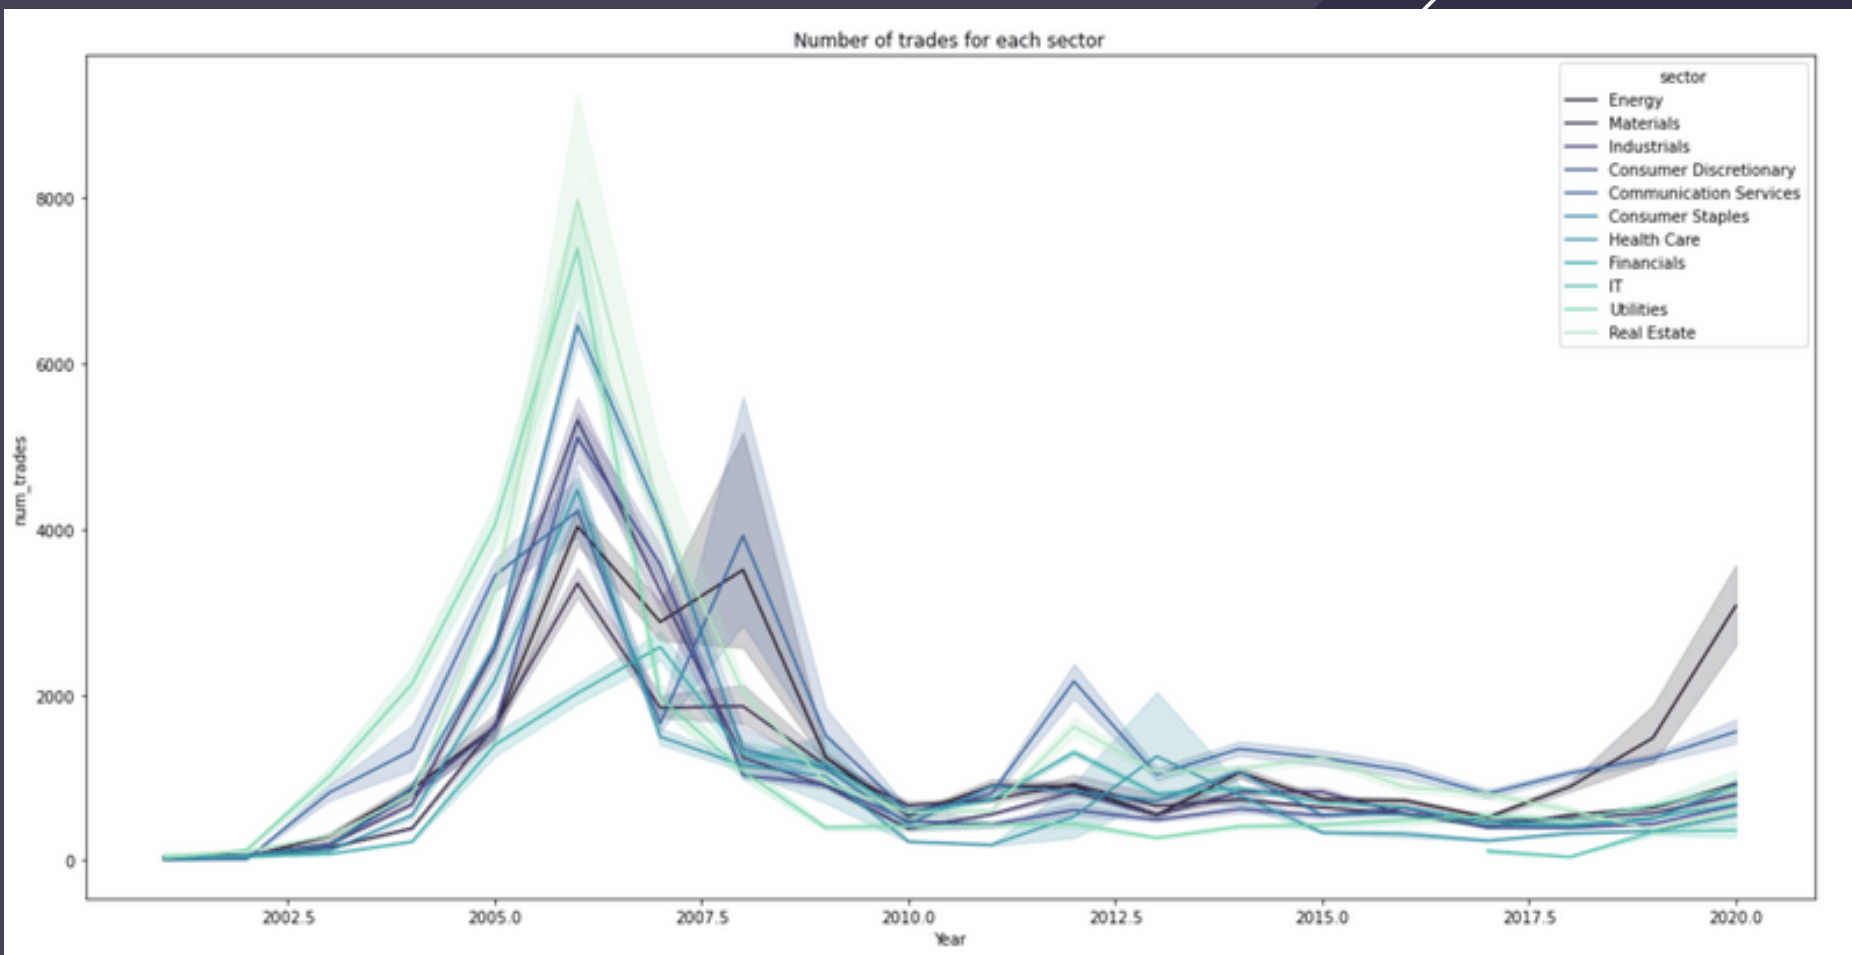


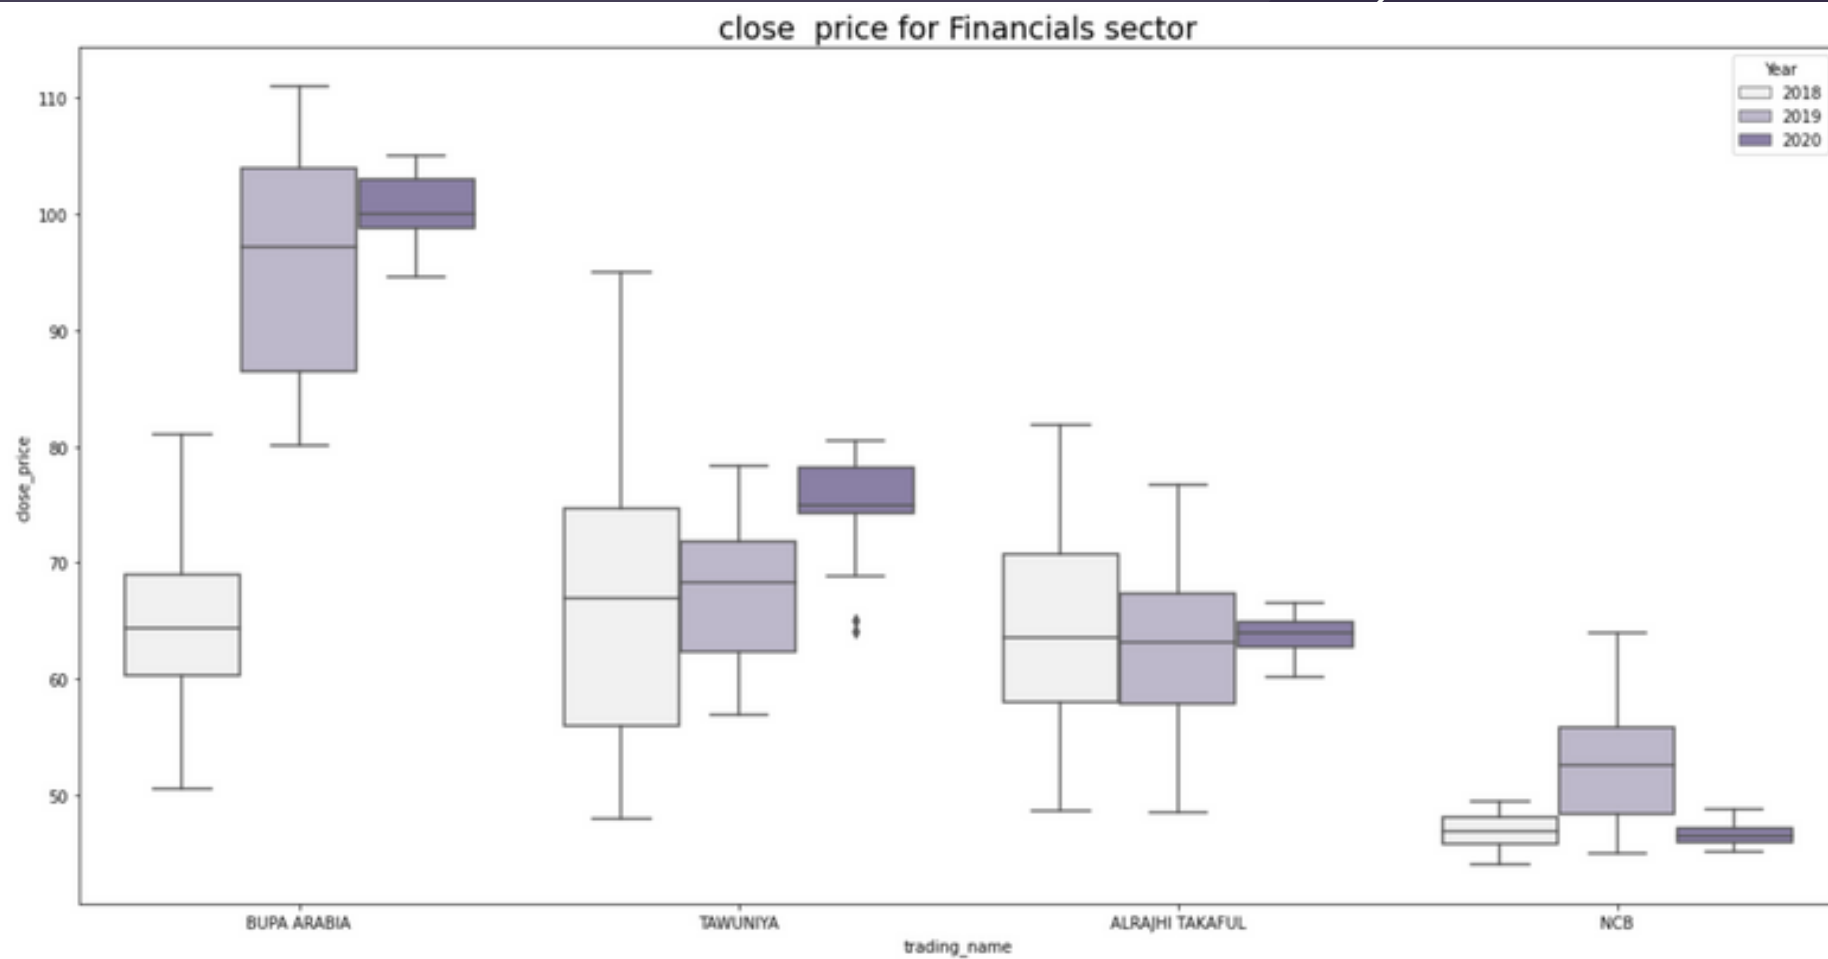


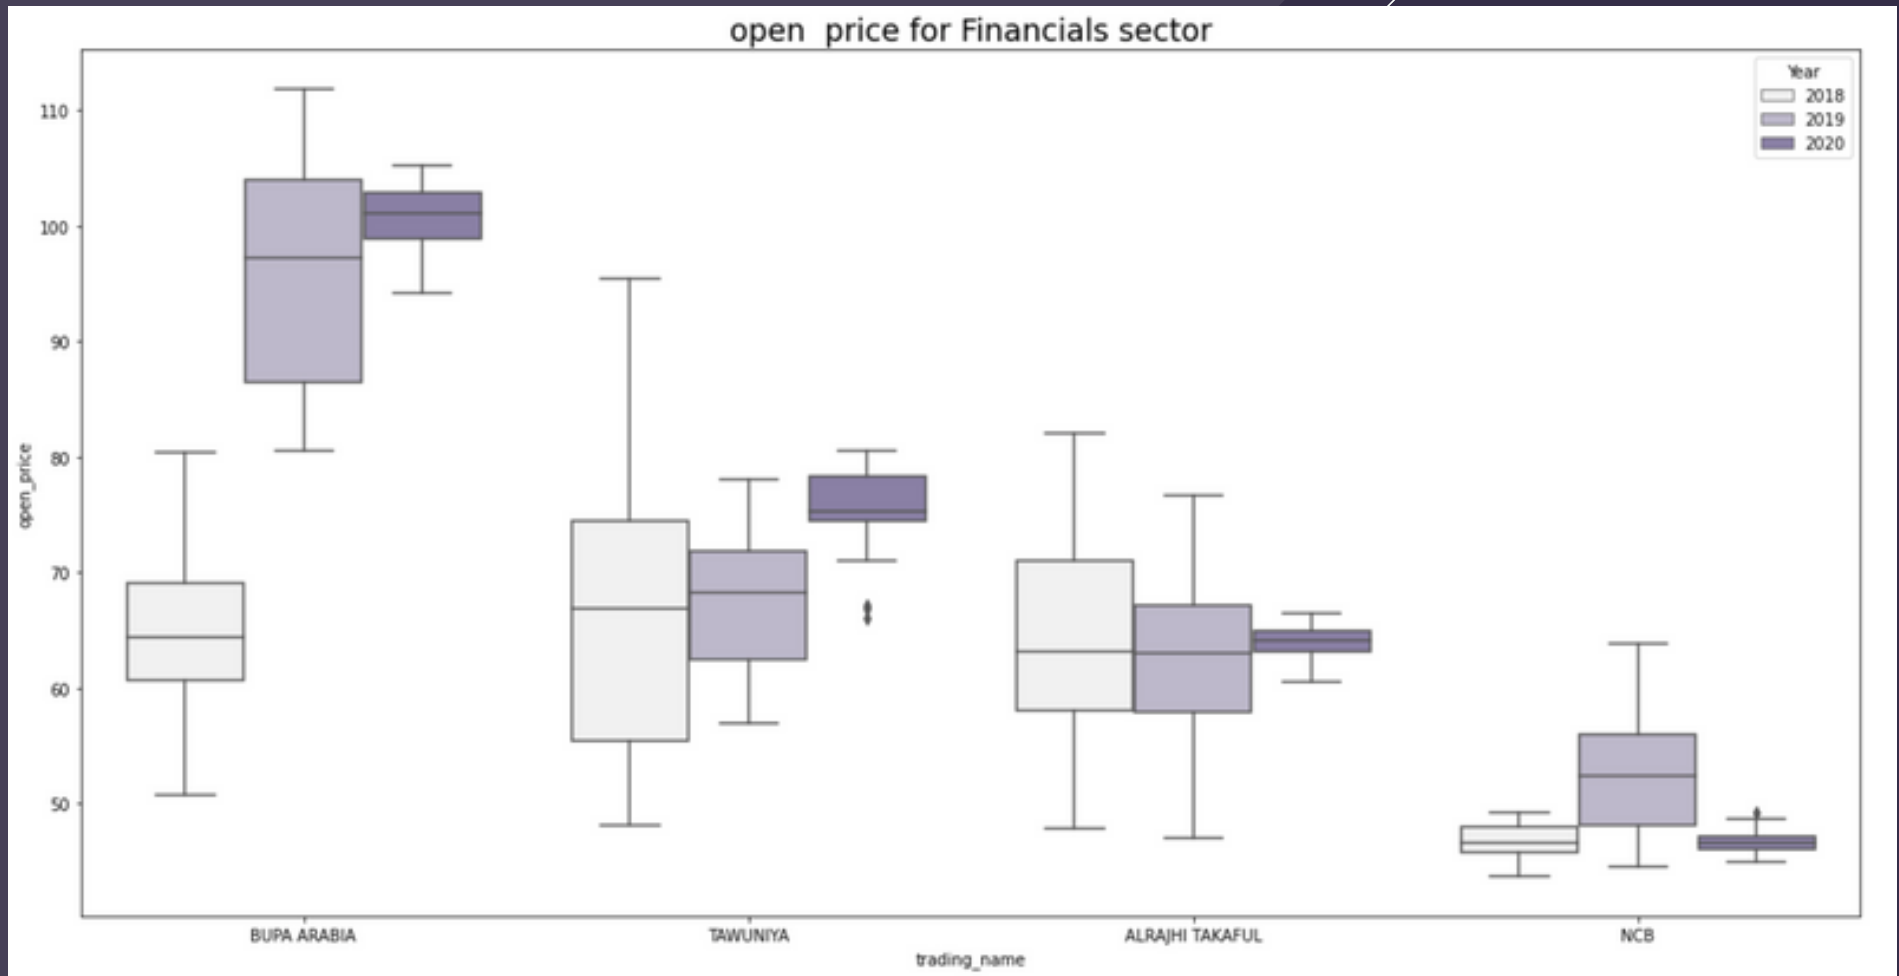


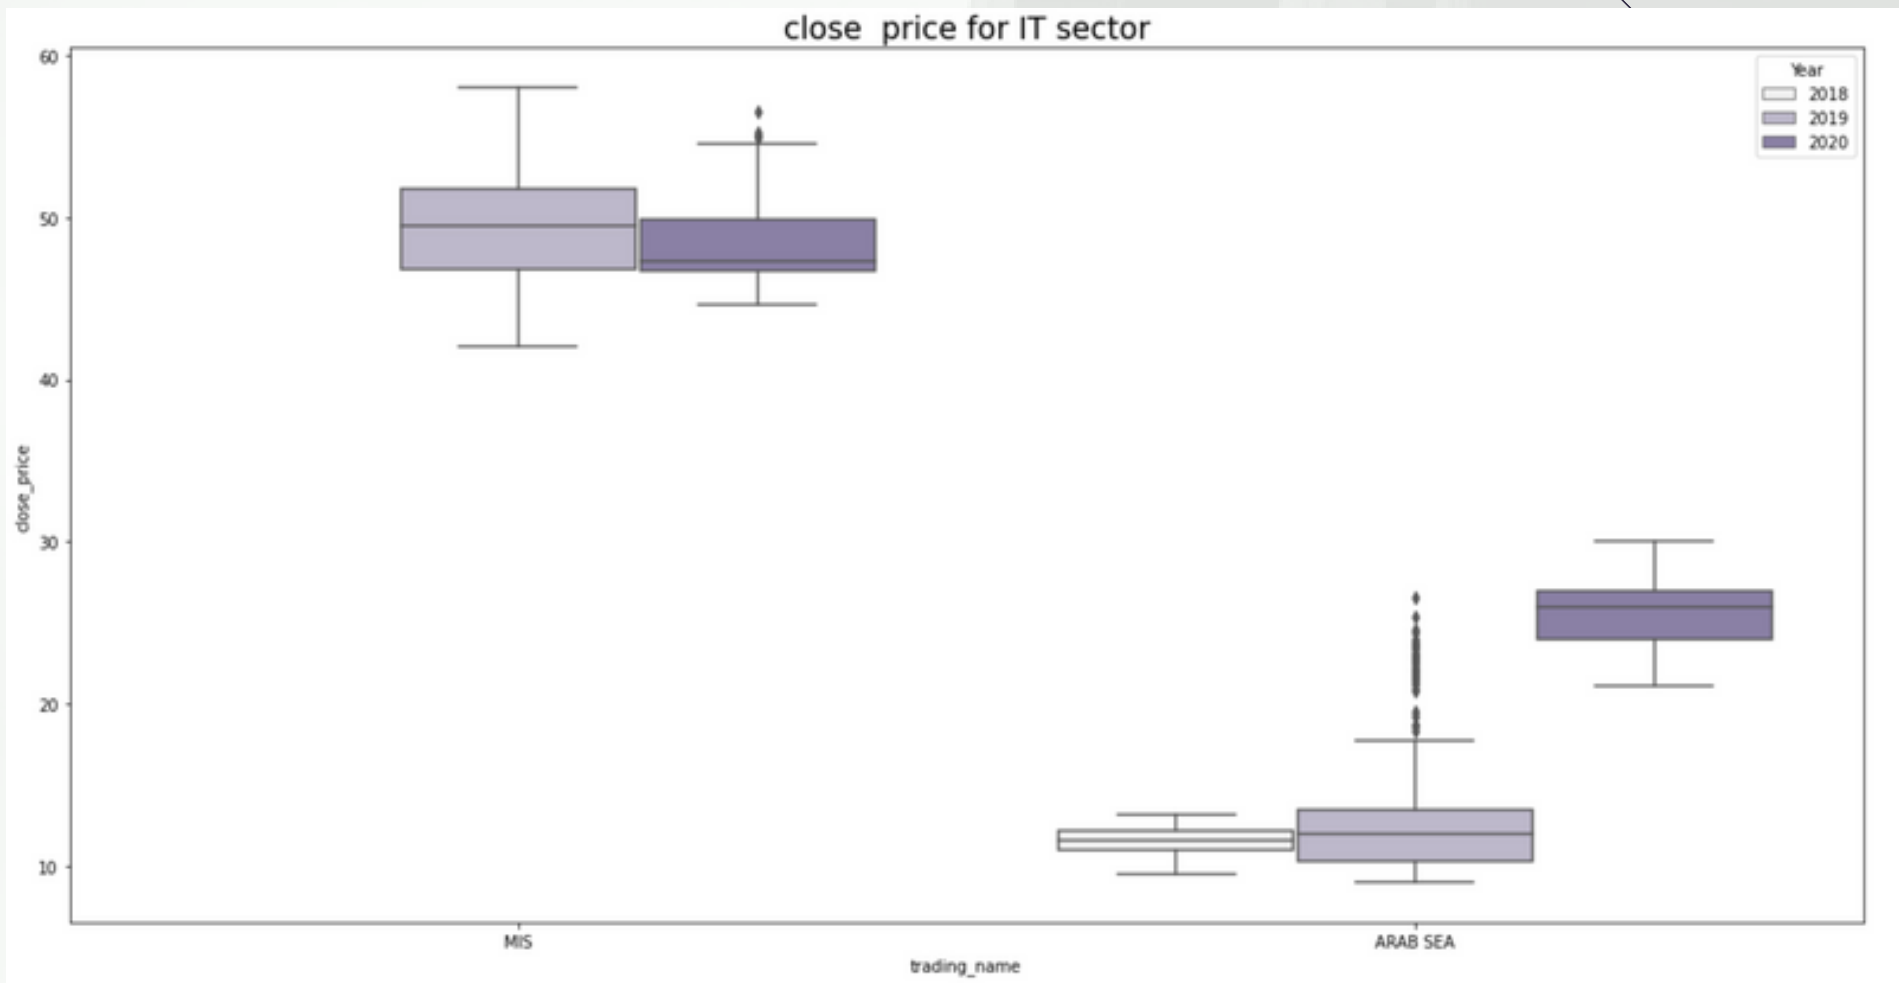


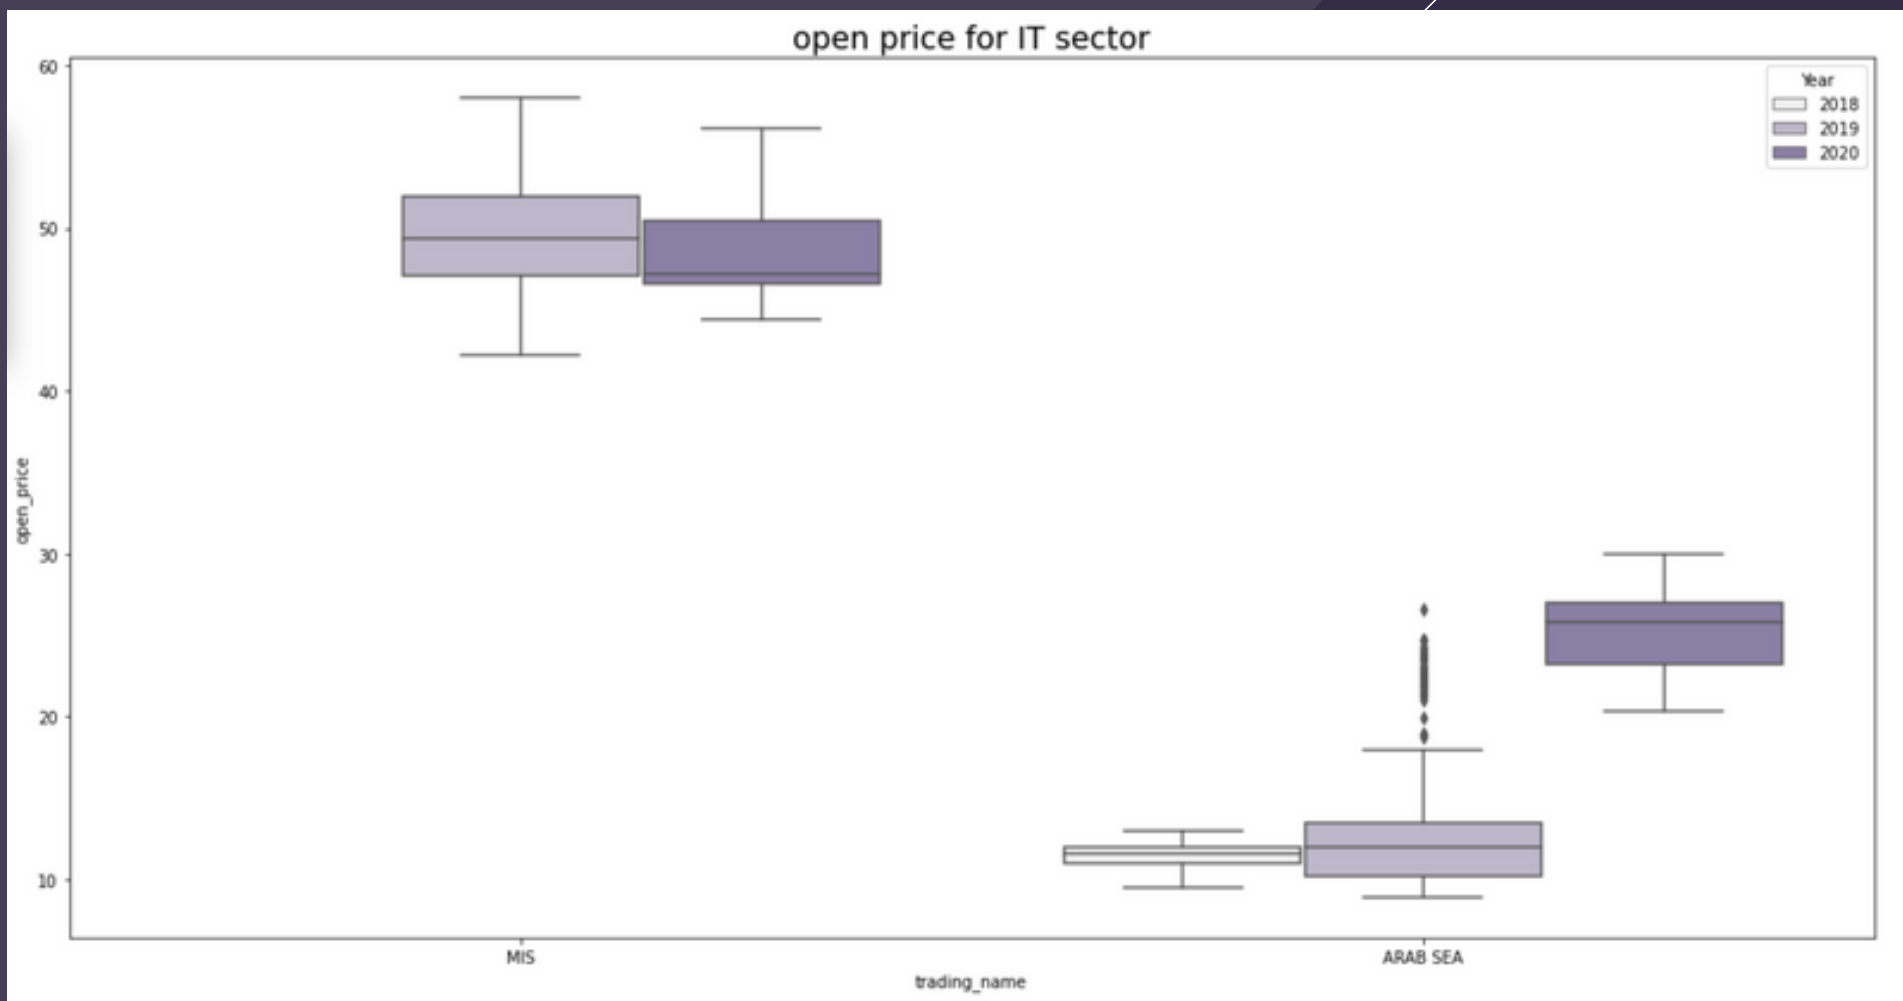








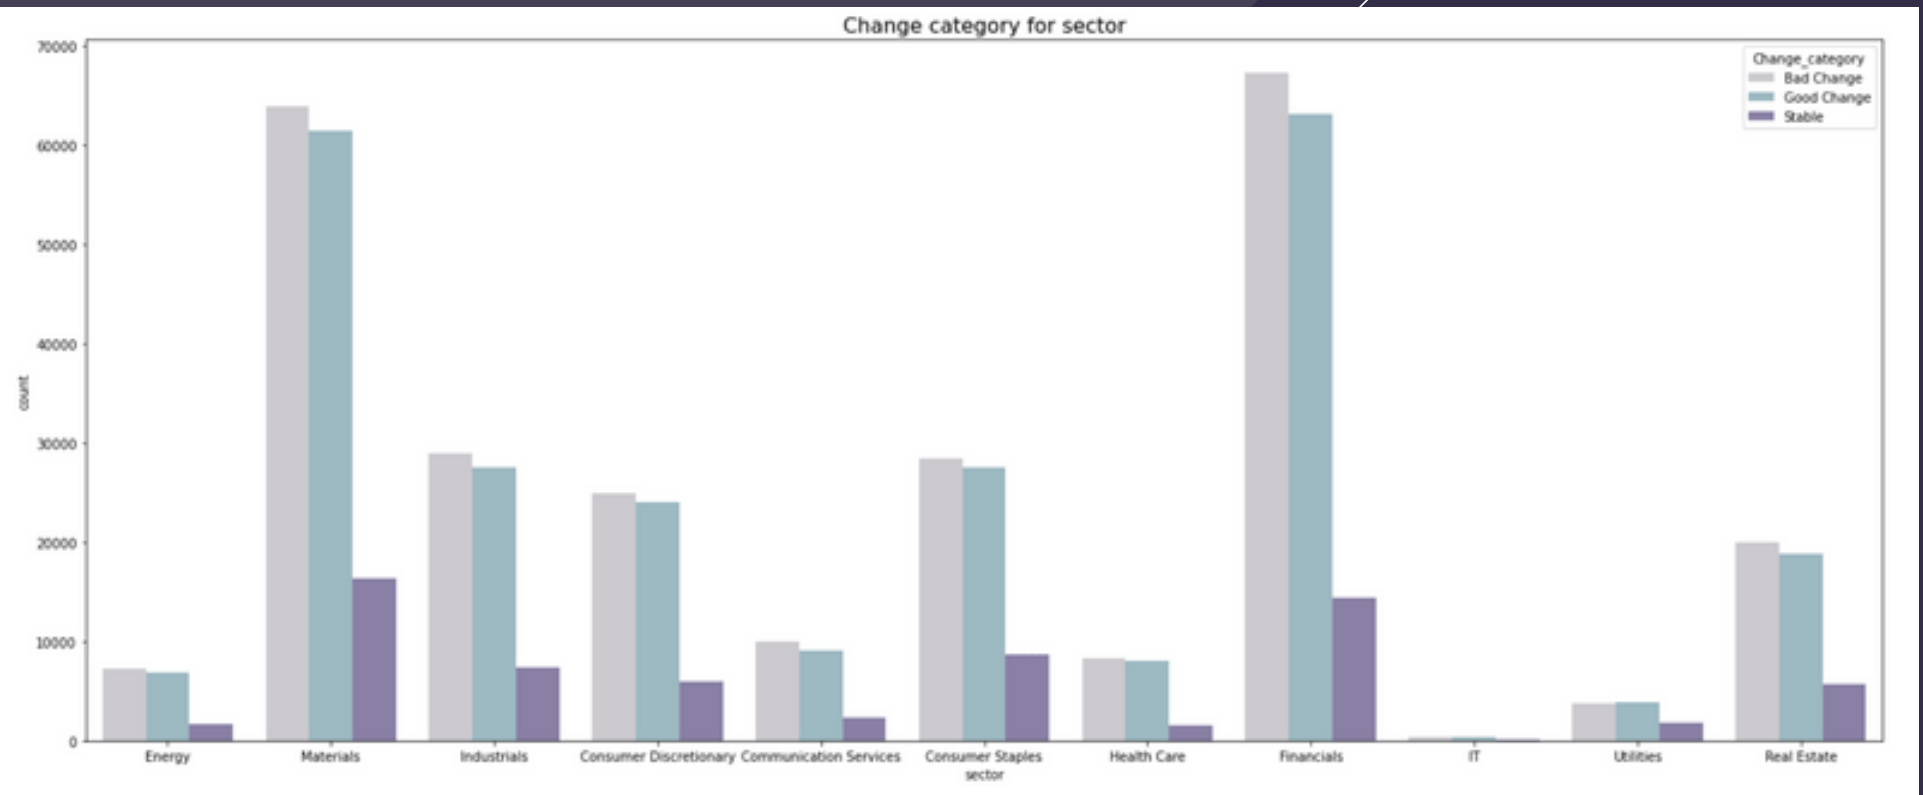


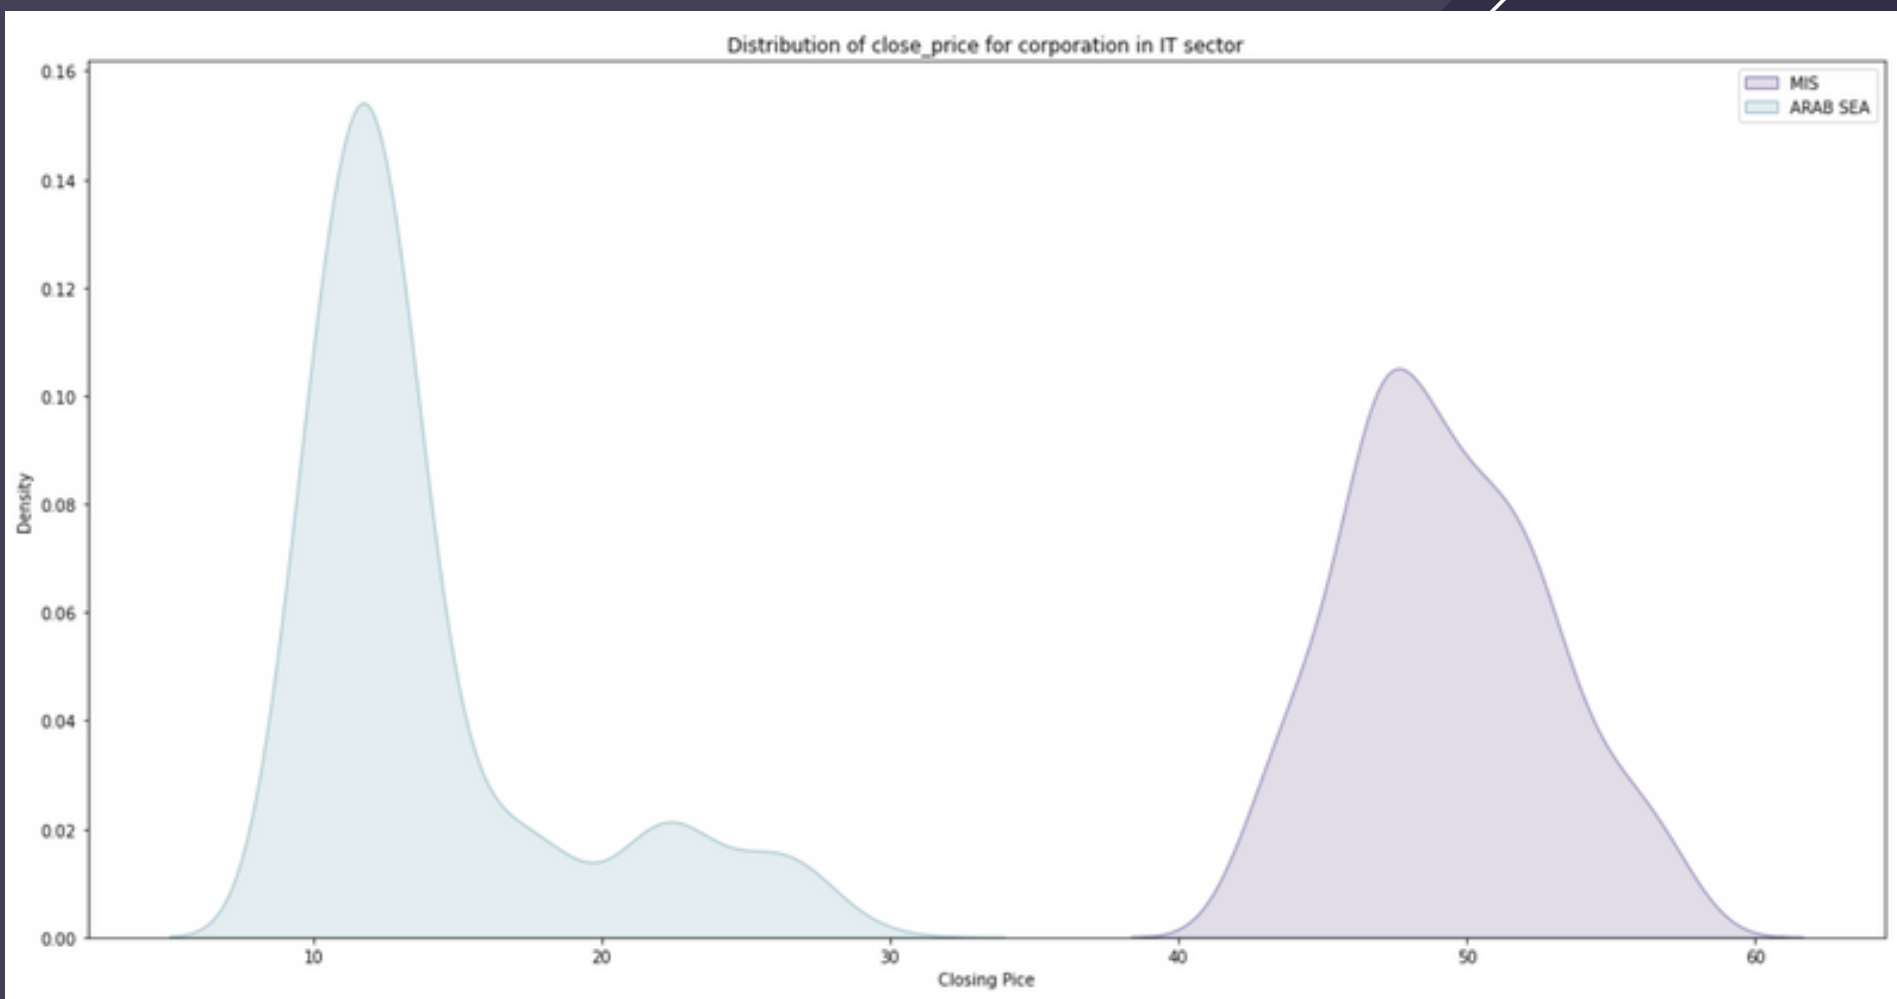


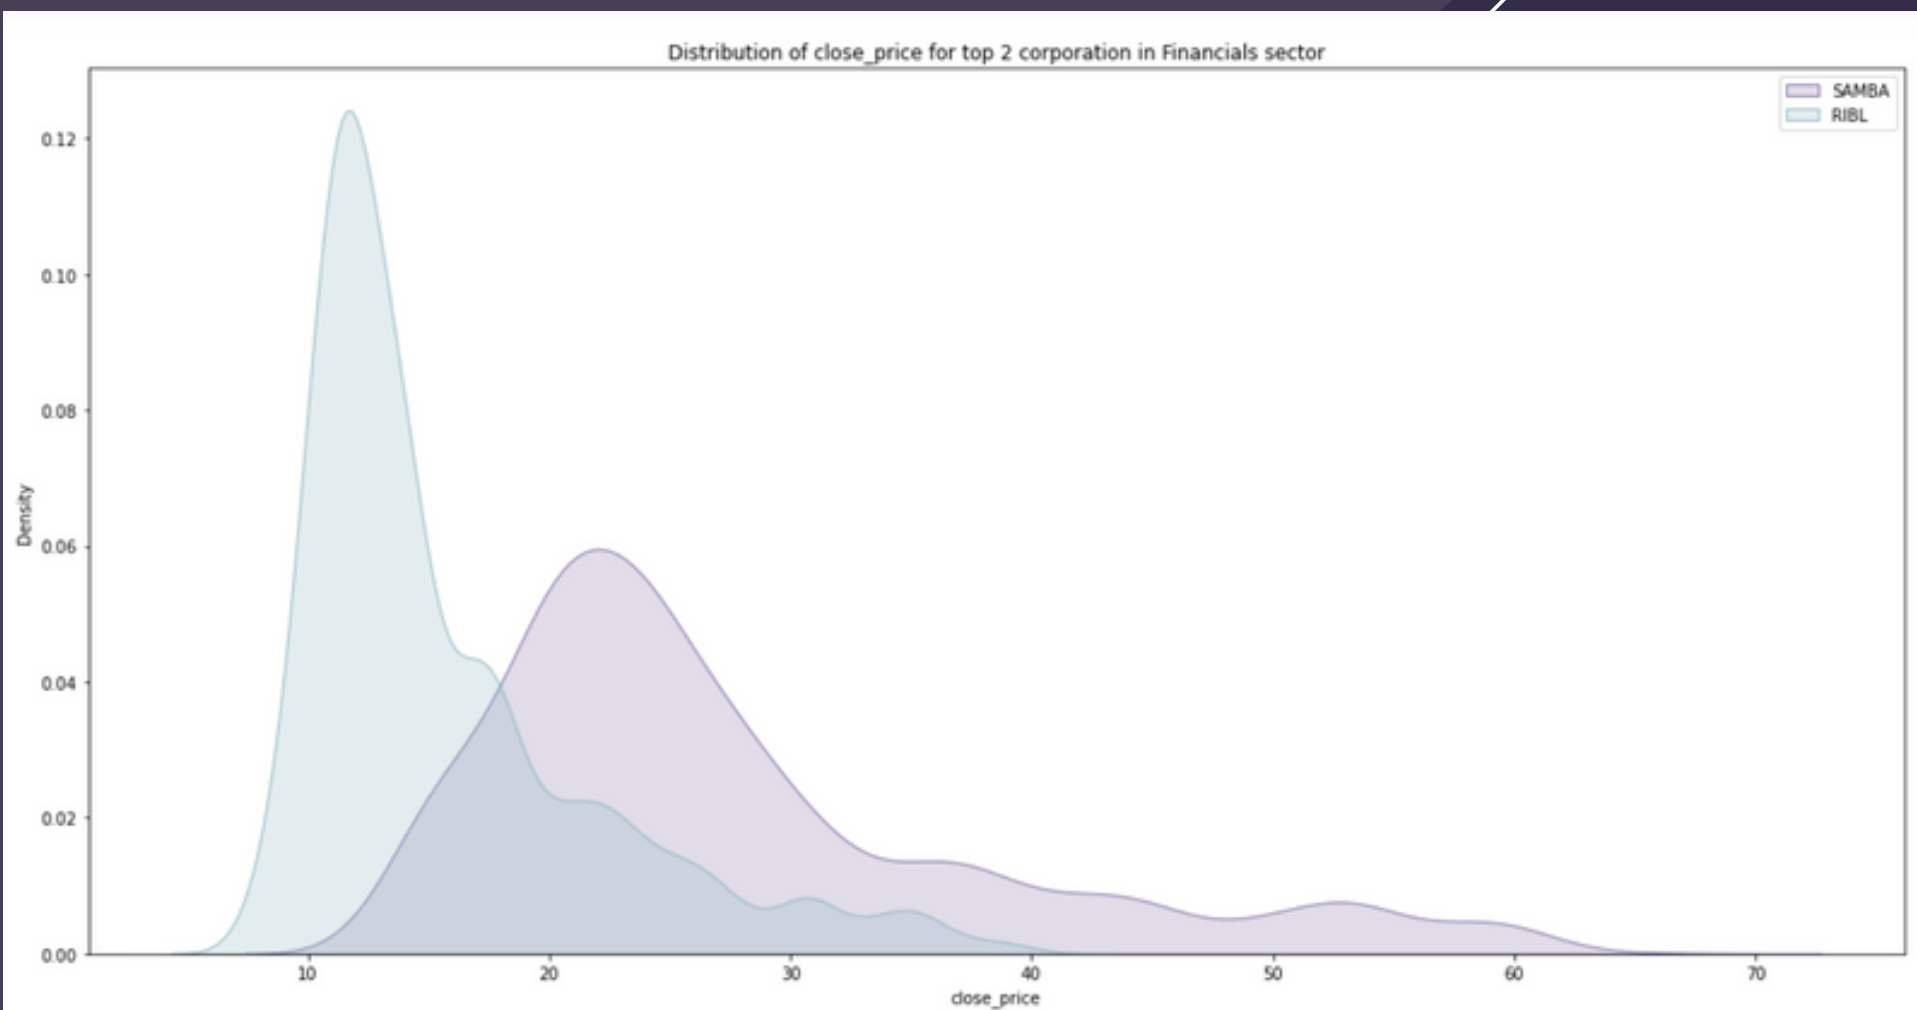





















Dashboards



Percentage of companies in each sector




The close price for each sectors from 2001 to 2020




Total change for each sector



Total Trades for each sector



Total of volume traded for each sector



Machine Learning Models



Regression Model



Classifications Model

SELECT THE FEATURE AND TARGET

```
target = 'close_price' # Target Variable  
features = ['open_price', 'low_price', 'change']
```

SCALER

```
scaler = StandardScaler()  
scaled_df = scaler.fit_transform(X)
```

SPLIT THE DATA

```
# split data into train and test  
# select random state = 42  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
```

BASELINE MODEL

```
# we will use the DummyRegressor model for the baseline
baseline_model = DummyRegressor()
baseline_model.fit(X_train,y_train)
baseline_model_pred = baseline_model.predict(X_test)

print(f"baseline model score: {r2_score(y_test, baseline_model_pred)}")
```

BASELINE MODEL RESULT

	Model	R2	MSE	MAE	RMSE
0	Baseline model	-0.000013	7217.246782	28.401061	84.954381

LINER REGRESSION

```
#create the model
liner = LinearRegression()
# fit the model using X train and y train
liner.fit(X_train , y_train)
# using X test to make our predication
linear_pred = liner.predict(X_test)
```

MODELS EVALUATION

```
# Create cost function that display all the cost functions for the regression models
def cost_function(pred):
    Adj_r2 = 1 - (1-r2_score(y_test, pred)) * (len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
    print("R Squared:",r2_score(y_test, pred))
    print("MSE:",mean_squared_error(y_test, pred))
    print("MAE:",mean_absolute_error(y_test, pred))
    print("RMSE:",np.sqrt(mean_squared_error(y_test, pred)))
    print("Adjusted R Squared:",Adj_r2)
```

LINER REGRESSION RESULT

	Model	R2	MSE	MAE	RMSE
1	Linear Regression	0.999606	2.844279	0.426416	1.686499

MODELS RESULTS

	Model	R2	MSE	MAE	RMSE
0	Baseline model	-0.000013	7217.246782	28.401061	84.954381
1	Linear Regression	0.999606	2.844279	0.426416	1.686499
2	Random Forest	0.987333	91.422137	5.789177	9.561492
3	KNN	0.999515	3.501239	0.301051	1.871160
4	GBR	0.999223	5.608102	0.671724	2.368143
5	XGB	0.999213	5.676292	0.396874	2.382497

MODEL OPTIMIZATION - HYPERPARAMETER TUNING

```
param_linear = {"fit_intercept": [True, False]}

random_linear_reg = RandomizedSearchCV(liner, # set the model
                                       param_linear, # set the parameter
                                       scoring='r2',
                                       verbose=1,
                                       n_jobs=-1)

# fit the model
random_linear_reg.fit(X_train, y_train)
# make the predication
random = random_linear_reg.predict(X_test)
```

LINER REGRESSION AFTER TUNING RESULT

	Model	R2	MSE	MAE	RMSE
2	Linear Regression with tuning	0.999606	2.844279	0.426416	1.686499

PIPELINE FOR BEST MODEL - LINER REGRESSION

```
numeric_features = X_train.describe().columns # Select the numrical feature

#Create Transformer for numerical data
numeric_transformer = Pipeline(
    steps=[
        ('imputer', SimpleImputer(strategy="most_frequent")),
        ('scaler', StandardScaler())
    ]
)

# Create a preprocessor transformer
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
    ]
)

# create the pipline
liner_reg_pp = Pipeline(
    steps=[
        ('preprocessor', preprocessor),# set the preprocessor

        ('reg',LinearRegression())# Create the Liner regression model
    ]
)
```

MODELS SUMMARY

- The best performance is the linear regression model followed by knn model .
- The worst performance is the Random forest regressor.
- Select the linear regression as the best model depend on the lowest RMSE value.



Machine Learning Models

Regression Model

Classifications Model

SELECT THE FEATURE AND TARGET

```
# set our features
cal_fatures= ['open_price', 'close_price', 'high_price', 'low_price']

#Set our target which is the the close price
cal_target= ['Change_category']

X= df[cal_fatures]
y= df[cal_target]
```

SCALER

```
scaler = StandardScaler()
scaled_df = scaler.fit_transform(X)
```

SPLIT THE DATA

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 42)
```

BASELINE MODEL

```
dummy1 = DummyClassifier()  
dummy1.fit(X_train, y_train)  
y_pred_dum = dummy1.predict(X_test)
```

BASELINE MODEL RESULT

	Model	Accuracy	Recall	precision	F1 score
0	Baseline model	0.454031	0.206144	0.454031	0.283548

SELECT THE FEATURE AND TARGET

```
# set our features
cal_fatures= ['open_price', 'close_price', 'high_price', 'low_price']

#Set our target which is the the close price
cal_target= ['Change_category']
```

RANDOM FOREST

```
# Create and fit the model
clas_forest = RandomForestClassifier(n_estimators = 100, criterion = 'gini', random_state = 42)
clas_forest.fit(X_train, y_train)
# Make orediction based on the test set
preds_ran = clas_forest.predict(X_test)
```

RANDOM FOREST RESULT

```
# Print the classification report
print(classification_report(y_test, preds_ran))
```

	precision	recall	f1-score	support
Bad Change	0.78	0.81	0.80	65770
Good Change	0.78	0.81	0.79	62655
Stable	0.49	0.36	0.41	16433
accuracy			0.76	144858
macro avg	0.69	0.66	0.67	144858
weighted avg	0.75	0.76	0.75	144858

MODELS RESULTS

	Model	Accuracy	Recall	precision	F1 score
0	Baseline model	0.454031	0.206144	0.454031	0.283548
1	Logistic Regression	0.758729	0.741616	0.758729	0.713941
2	Random Forest	0.758529	0.749616	0.758529	0.752706
3	KNN	0.756320	0.744912	0.756320	0.744185
4	GBC	0.693741	0.694983	0.693741	0.666824

MODEL OPTIMIZATION - HYPERPARAMETER TUNING

```
parameters = {'bootstrap': [True, False],
              'max_depth': [10, 20, 30, None],
              'max_features': ['auto', 'sqrt'],
              'min_samples_leaf': [1, 2, 4],
              'min_samples_split': [2, 5, 10],
              'n_estimators': [50, 80, 100]}

random_search = RandomizedSearchCV(clas_forest, # Model
                                  parameters, # Parameters to tune
                                  cv=5, # Cross Validation
                                  verbose=1, # Shows output while training
                                  n_jobs=-1, # How many core to use on your computer (-1 means use all cores)
                                  scoring="accuracy"
                                  )

# Fit the model
random_search.fit(X_train, y_train)
```

LINER REGRESSION AFTER TUNING RESULT

	Model	Accuracy	Recall	precision	F1 score
2	Random Forest with tuning	0.750956	0.758492	0.768822	0.760486

PIPELINE FOR BEST MODEL - RANDOM FOREST

```
numeric_features = X_train.describe().columns

# Create a transformer for numeric columns
numeric_transformer = Pipeline(
    steps=[
        # missing values --> by default mean
        ('scaler', StandardScaler())
    ]
)

Random_regression = Pipeline(
    steps=[
        ('Randomregression', RandomForestClassifier(n_estimators = 250, max_depth = 50, random_state = 42))
    ]
)

# Fit the model
Random_regression.fit(X_train, y_train)
```



MODELS SUMMARY

We have tried several classification models to come up with the best model that can predict the change category depending on the high price, low price, open price, and close price.

Then we select the Random forest as the best model based on the highest weighted avg of f1.



RESULTS



Future Work

Thank You For Listening

ANY QUESTION?

TEAM MEMBERS :



Maha



Fatima



Rasha



Joharah



Areej



Samar

